

SEV-Cap

Grounded Multi-Style Video Captioning

Four captions per clip — formal, sarcastic, humorous-tech, humorous-non-tech — optimized for the same two axes the LLM judge grades: factual accuracy and stylistic tone.

accuracy

tone

anytime I/O

linux/amd64

Single-pass captioners fail the judge

Track 2 scores each caption on accuracy (1–5) and tone (1–5). Combined leaderboard score is (mean accuracy + mean tone) / 10.

Accuracy fails

Models invent cities, brands, camera moves, and plot that never appear in the frames.

Tone fails

A “sarcastic” caption that is just a formal paraphrase, or humor with no joke.

Harness fails

Wrong schema, missing tasks, or timeouts → OUTPUT_MISSING / INVALID_RESULTS.

Optimize for how the judge actually scores

01 One grounded description

Describe the scene from keyframes, then self-verify against the frames before any styled writing.

02 Multi-candidate styles

Draft several captions per style (text + vision-grounded). Never ship the first draft blindly.


03 Vision prejudice

Score each candidate on accuracy + tone — the same axes as Track 2 — and keep the best.

04 Polish weak styles

If a style is still soft, reselect with safer / vision-grounded rewrites before time runs out.

End-to-end scoring pipeline



- Shared description anchors every style so jokes stay on the same facts.
- Prejudge failures retry with JSON/regex parsing — no fake mid-scores that hide bad picks.
- Anytime writes: placeholders early, atomic upgrades, finalize every task_id.

Four styles — same facts, different voice

Formal

Precise, neutral, news-desk prose. Concrete visible details only.

Sarcastic

Dry ironic praise or understatement. Tone without invented plot.

Humorous · Tech

One tech metaphor mapped onto the primary on-screen action.

Humorous · Non-Tech

Warm everyday personification. Zero jargon, no invented drama.

What the scoring harness runs

Input	/input/tasks.json → [{task_id, video_url, styles}]
Output	/output/results.json → [{task_id, captions{...}}]
Image	ghcr.io/skx56/sevcap-grounded:latest
Platform	linux/amd64 · argument-free ENTRYPOINT · FIREWORKS_API_KEY at runtime
Defaults	SEVCAP_CANDIDATES=3 · SEVCAP_POLISH=1 · SEVCAP_AUDIO=0 · Kimi K2.6

Placeholders are written immediately; every expected task_id is finalized. Timeouts may degrade quality but must not produce MISSING_TASKS or OUTPUT_MISSING.

Official sample-style clips

0.966

combined (acc+tone)/10

mean acc 4.75 · tone 4.91

8 / 8

tasks sev-verified

Harness OK

schema + task coverage

Docker OK

ghcr.io/skx56/sevcap-grounded

Internal vision judge mirrors Track 2 axes on the eight AMD sample clips (boulevard, kitten, office, mountains, waves, intersection, cooking, track).

Stack & links

- Kimi K2.6 on Fireworks (VLM + text)
- Python 3.11 · asyncio · ffmpeg
- Docker linux/amd64 → GHCR
- Optional Gemma override for bonus chase

Submit this image

ghcr.io/skx56/sevcap-grounded:latest

Repo: github.com/skx56/sev-cap Key: FIREWORKS_API_KEY
only at runtime Never baked into the image